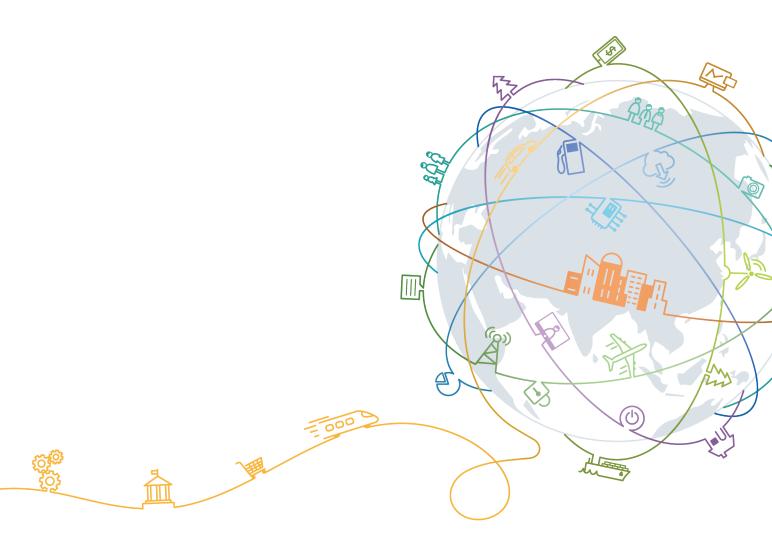
OpenStack Cinder Driver

Configuration Guide

Issue 01

Date 2021-02-01





Copyright © Huawei Technologies Co., Ltd. 2021. All rights reserved.

No part of this document may be reproduced or transmitted in any form or by any means without prior written consent of Huawei Technologies Co., Ltd.

Trademarks and Permissions

HUAWEI and other Huawei trademarks are trademarks of Huawei Technologies Co., Ltd. All other trademarks and trade names mentioned in this document are the property of their respective holders.

Notice

The purchased products, services and features are stipulated by the contract made between Huawei and the customer. All or part of the products, services and features described in this document may not be within the purchase scope or the usage scope. Unless otherwise specified in the contract, all statements, information, and recommendations in this document are provided "AS IS" without warranties, guarantees or representations of any kind, either express or implied.

The information in this document is subject to change without notice. Every effort has been made in the preparation of this document to ensure accuracy of the contents, but all statements, information, and recommendations in this document do not constitute a warranty of any kind, express or implied.

Huawei Technologies Co., Ltd.

Address: Huawei Industrial Base

Bantian, Longgang Shenzhen 518129

People's Republic of China

Website: https://www.huawei.com

Email: support@huawei.com

Contents

1 Overview	1
2 Version Mappings	2
2.1 Version Mappings Among Cinder Driver, Huawei Storage, and OpenStack	2
2.2 Mappings Among Cinder Driver, Features, and OpenStack Versions	3
2.3 Mappings Between Cinder Driver and Storage Licenses	5
3 Deployment	7
3.1 Obtaining Cinder Driver	
3.2 Deploying Cinder Driver for the Cinder Volume Service in Non-Containerized Mode	8
3.3 Deploying Cinder Driver for the Cinder Volume Service in Containerized Mode	9
4 Configuring Basic Properties	10
4.1 Configuring Cinder Driver for iSCSI-based Storage	10
4.2 Configuring Cinder Driver for FC-based Storage	14
4.3 Configuring Cinder Driver in Container-based Deployment	17
5 Configuring the Multipathing Software	18
5.1 Configuring Host Multipathing	18
5.1.1 Installing the Multipathing Tool Package	18
5.1.2 Configuring the Multipathing Service	19
5.2 Configuring OpenStack Multipathing	19
5.2.1 Configuring Multipathing for OpenStack Nova Nodes	19
5.2.2 Configuring Multipathing for OpenStack Cinder Nodes	20
6 Configuring Advanced Properties	21
6.1 Configuring the Thick Property	22
6.2 Configuring the Thin Property	22
6.3 Configuring the SmartQoS Property	23
6.4 Configuring the SmartPartition Property	24
6.5 Configuring the SmartCache Property	24
6.6 Configuring the SmartTier Property	25
6.7 Configuring the HyperMetro Property	25
6.8 Configuring the Remote Replication Property	27
6.9 Creating a LUN in a Specified Storage Pool	28
6.10 Creating a LUN on a Disk of a Specified Type	29

6.11 Creating a LUN in on a Specified Backend	29
6.12 Configuring a Backup Snapshot	. 30
6.13 Configuring Certificate Authentication	30
6.14 Configuring the Owning Controller of a LUN	. 30
6.15 Configuring the SmartDedupe Property	31
6.16 Configuring the SmartCompression Property	32
6.17 Creating a Volume Using fastclone	
6.18 Configuring the LUN Copy Speed	33
6.19 Configuring the Application Type of a LUN	
6.20 Configuring a Tenant User	
6.21 Configuring Multiple Backends	. 35
6.22 Configuring Auto Zoning	
6.23 Configuring CHAP	. 35
7 Best Practices	37
7.1 Configuring Basic Properties for Interconnection Between iSCSI and OceanStor Dorado V3	. 37
7.2 Initiator Configuration Policy Example	38
8 FAQs	40
8.1 Do I Need to Create Hosts, Host Groups, LUN Groups, and Other Associated Objects on Huawei Storage Before Using Cinder Driver to Map Volumes?	40
8.2 What Can I Do if Mapping Fails Because Nova Logs Indicate that the lsblk Command Fails to Be Executed During Volume Mapping?	40
8.3 Create Volume From Image Failed When verify_glance_signatures=enabled	41
8.4 NewNode1	41
8.5 NewNode2	41
9 Appendix	.42
9.1 Appendix 1: ALUA Parameter Description	42
9.2 Appendix 2: Configuring Volume-backend image	43
9.3. Appendix 3: Configuring Image-Volume Cache	11

1 Overview

Huawei Cinder Driver is a plug-in that is deployed on the OpenStack Cinder module. The plug-in connects Huawei storage to provide the volume storage capability for the OpenStack platform. Cinder Driver supports iSCSI and Fibre Channel (FC) protocols.

2 Version Mappings

This chapter describes the version mappings among Cinder Driver, Huawei storage, and OpenStack as well as mappings between Cinder Driver and storage licenses.

- 2.1 Version Mappings Among Cinder Driver, Huawei Storage, and OpenStack
- 2.2 Mappings Among Cinder Driver, Features, and OpenStack Versions
- 2.3 Mappings Between Cinder Driver and Storage Licenses

2.1 Version Mappings Among Cinder Driver, Huawei Storage, and OpenStack

Table 2-1 Version mappings among Cinder Driver, Huawei storage, and OpenStack

OpenStack	Huawei Storage
• Queens	OceanStor Dorado V6 6.0/6.1
Rocky	OceanStorDorado V3 V300R001/V300R002
• Stein	OceanStor F V5/V5 V500R007/V500R007 KunPeng
• Train	OceanStor F V3/V3 V300R003/V300R006
Ussuri	
• Victoria	

□ NOTE

Huawei only maintains the drivers for the latest 6 stable OpenStack versions. Though the earlier versions can still be retrieved from Huawei warehouse website, they are out of maintenance. Be wary of using them.

2.2 Mappings Among Cinder Driver, Features, and OpenStack Versions

Table 2-2 Mappings among Cinder Driver, features, and OpenStack versions (√: supported, x: not supported)

Feature	Normal LUN	HyperMetro LUN	Remarks
OpenStack	Queens/ Rocky/Stein/ Train/Ussuri/ Victoria	Queens/Rocky/ Stein/ Train/Ussuri/ Victoria	The following features are applicable to the scenario where HyperMetro storage is normal.
Create Volume	√	√	N/A
Delete Volume	√	√	N/A
Attach Volume	√	√	N/A
Detach Volume	√	√	N/A
Extend Volume	√	√	N/A
Create Snapshot	√	√	HyperMetro LUN supports snapshot creation only on the local end of HyperMetro.
Delete Snapshot	√	√	HyperMetro LUN supports snapshot deletion only on the local end of HyperMetro.
Create Volume from Snapshot	√	√	HyperMetro LUN supports snapshot-based volume creation only on the local end of HyperMetro.
Create Volume from Image	√	√	N/A
Create Volume from Volume	√	√	N/A
Create Image from Volume	√	√	N/A
Volume Migration	√	х	N/A
Thin Volume	√	√	N/A
Thick Volume	√	√	N/A

Feature	Normal LUN	HyperMetro LUN	Remarks
QoS	√	√	HyperMetro LUN supports SmartQoS configuration only on the local end of HyperMetro.
FC zoning	√	√	N/A
Manage/Unmanage Volume	√	х	N/A
Manage/Unmanage Snapshot	√	х	N/A
Replication V2.1	√	х	N/A
Backup Snapshot	√	√	HyperMetro LUN supports snapshots backup only on the local end of HyperMetro.
Retype	√	х	N/A
Multipath	√	√	N/A
Multi-Attach	√	√	N/A
Consistency Group	√	х	N/A
Snapshot Consistency Group	√	√	HyperMetro LUN supports snapshot creation only on the local end of HyperMetro.
HyperMetro Consistency Group	x	√	N/A
Replication Consistency Group	√	х	N/A
SmartTier	√	√	HyperMetro LUN supports SmartTier configuration only on the local end of HyperMetro.
SmartCache	√	√	HyperMetro LUN supports SmartCache configuration only on the local end of HyperMetro.
SmartPartition	√	√	HyperMetro LUN supports SmartPartition configuration only on the local end of HyperMetro.
Create a volume on a certain storage pool	√	х	N/A
Create a volume with certain disk type	√	Х	N/A

NOTICE

- OceanStor Dorado does not support Thick Volume, SmartTier, SmartCache, and SmartPartition.
- Entry-level storage does not support some value-added features. For details, refer to the specifications of the corresponding storage model.

2.3 Mappings Between Cinder Driver and Storage Licenses

Table 2-3 Mappings between Cinder Driver and storage licenses

Feature	Storage License
Create Volume	N/A
Delete Volume	N/A
Attach Volume	N/A
Detach Volume	N/A
Extend Volume	N/A
Create Snapshot	HyperSnap
Delete Snapshot	N/A
Create Volume from Snapshot	HyperCopy/HyperClone
Create Volume from Image	HyperSnap&(HyperCopy/HyperClone)
Create Volume from Volume	HyperSnap&(HyperCopy/HyperClone)
Create Image from Volume	N/A
Volume Migration	SmartMigration
QoS	SmartQoS
FC zoning	N/A
SmartTier	SmartTier
SmartCache	SmartCache
Thin Volume	SmartThin
Thick Volume	N/A
SmartPartition	SmartPartition
HyperMetro	HyperMetro
Retype	N/A

Feature	Storage License
Manage/Unmanage Volume	N/A
Manage/Unmanage Snapshot	N/A
Replication V2.1	HyperReplication
Create a volume on a certain storage pool	N/A
Create a volume with certain disk type	N/A
Backup Snapshot	HyperCopy/HyperClone
Multipath	N/A
Multi-Attach	N/A
Consistency Group	N/A
Snapshot Consistency Group	HyperSnap
HyperMetro Consistency Group	HyperMetro
Replication Consistency Group	HyperReplication

3 Deployment

- 3.1 Obtaining Cinder Driver
- 3.2 Deploying Cinder Driver for the Cinder Volume Service in Non-Containerized Mode
- 3.3 Deploying Cinder Driver for the Cinder Volume Service in Containerized Mode

3.1 Obtaining Cinder Driver

You can obtain Cinder Driver in either of the following ways:

- OpenStack community warehouse. Since the Kilo version, Cinder Driver has been integrated into the OpenStack community warehouse. After OpenStack is installed, Cinder Driver is provided. Cinder Driver is stored in the ../cinder/ cinder/volume/drivers/huawei directory.
- Huawei OpenStack Driver warehouse
- **Step 1** Open a browser and enter https://github.com/Huawei/OpenStack_Driver in the address box.
- **Step 2** Click the **Download ZIP** button to download the Cinder Driver package to a local host.
- **Step 3** Decompress the package.
- **Step 4** In the decompressed directory, find the **Cinder** directory, which contains Cinder Driver of multiple OpenStack versions.

□ NOTE

- After OpenStack community versions are released, new features cannot be integrated.
 The bug fixing period is long and risks exist. Huawei OpenStack Driver warehouse
 versions ensure that new features are released and bugs are fixed in time. These
 versions are more stable than OpenStack community versions.
- The OpenStack community maintains only two stable versions. Huawei OpenStack
 Driver warehouse maintains six stable versions, ensuring long-term stable running of
 historical versions.
- It is strongly recommended that OpenStack community versions be replaced by Huawei OpenStack Driver warehouse versions.

3.2 Deploying Cinder Driver for the Cinder Volume Service in Non-Containerized Mode

Step 1 Search for the original Huawei Driver code in the system. The relative path is .../ cinder/volume/drivers/huawei.

○ NOTE

The absolute path may vary in different systems. You can obtain the Cinder Driver code directory by running the following commands:

```
# python -c "from cinder.volume.drivers import huawei; print huawei.__path__" ['/usr/lib/python2.7/site-packages/cinder/volume/drivers/huawei']
```

The driver path is:

/usr/lib/python2.7/site-packages/cinder/volume/drivers/huawei/

- **Step 2** Delete all Huawei Driver code files in the directory returned in step 1. Copy obtained Cinder Driver code files to the directory.
- **Step 3** Copy all Huawei Driver code files of the corresponding version to the directory returned in **Step 1**.

```
# ls -l
-rw-r--r-- 1 root root 4787 Mar 4 23:40 constants.py
-rw-r--r-- 1 root root 10316 Mar 4 23:40 fc_zone_helper.py
-rw-r--r-- 1 root root 17905 Mar 4 23:40 huawei_conf.py
-rw-r--r-- 1 root root 125657 Mar 4 23:48 huawei_driver.py
-rw-r--r-- 1 root root 24276 Mar 4 23:40 huawei_t.py
-rw-r--r-- 1 root root 6698 Mar 4 23:40 huawei_utils.py
-rw-r--r-- 1 root root 18437 Mar 4 23:40 hypermetro.py
-rw-r--r-- 1 root root 38948 Mar 4 23:40 __init__.py
-rw-r--r-- 1 root root 97236 Mar 4 23:40 rest_client.py
-rw-r--r-- 1 root root 8834 Mar 4 23:40 smartx.py
```

Or

```
# ls -l
-rw-r--r-- 1 root root 5631 Mar 4 23:40 constants.py
-rw-r--r-- 1 root root 33146 Mar 4 23:40 huawei_base_driver.py
-rw-r--r-- 1 root root 21774 Mar 4 23:40 huawei_conf.py
-rw-r--r-- 1 root root 11512 Mar 4 23:40 huawei_driver.py
-rw-r--r-- 1 root root 98370 Mar 4 23:40 huawei_flow.py
-rw-r--r-- 1 root root 19599 Mar 4 23:40 huawei_utils.py
-rw-r--r-- 1 root root 14733 Mar 4 23:40 hypermetro.py
-rw-r--r-- 1 root root 21 Mar 4 23:40 __init__.py
-rw-r--r-- 1 root root 20056 Mar 4 23:40 replication.py
-rw-r--r-- 1 root root 4938 Mar 4 23:40 smartx.py
```

3.3 Deploying Cinder Driver for the Cinder Volume Service in Containerized Mode

Step 1	Log in to the node with the Cinder Volume container image. Run the docker save
	command to back up the Cinder Volume container image.

NOTE

To obtain the image version, run the docker image ls command.

Step 2 Run the following command to create temporary directory **build** in any directory and go to the created directory:

mkdir build;cd build

Step 3 Run the following commands to copy the obtained Cinder Driver code files to the current directory:

Step 4 Create the **Dockerfile** file in the current directory and edit the following content:

FROM ***:***
COPY *.py /huawei/cinder/driver/path/

□ NOTE

- Replace ********* with the name and version of the original Cinder Volume container image.
- Replace /huawei/cinder/driver/path/ with the path of Cinder Driver in the container.
 You can obtain the Cinder Driver code directory by running the following commands:

python -c "from cinder.volume.drivers import huawei; print huawei.__path__" ['/usr/lib/python2.7/site-packages/cinder/volume/drivers/huawei']

Step 5 Run the following command to build an image:

docker build -t ***:*** .

□ NOTE

Replace ***:*** with the name and version of the original Cinder Volume container image.

4 Configuring Basic Properties

This chapter describes how to configure Huawei Cinder Driver for iSCSI-and FC-based storage.

NOTICE

- Ensure that the storage pool used for configuring Cinder Driver exists on Huawei storage. Otherwise, create a storage pool. In addition, the type of this storage pool must be block storage service.
- For an FC network, the sg tool must be installed so that storage resources can be used.
- All parameter values in the Cinder Driver configuration file cannot contain the following special characters: <>&'''
- Configuring Cinder Driver in container-based deployment is different from that in non-container-based deployment. For details, see 4.3 Configuring Cinder Driver in Container-based Deployment.
- 4.1 Configuring Cinder Driver for iSCSI-based Storage
- 4.2 Configuring Cinder Driver for FC-based Storage
- 4.3 Configuring Cinder Driver in Container-based Deployment

4.1 Configuring Cinder Driver for iSCSI-based Storage

This section describes how to configure Cinder Driver for iSCSI-based storage products.

- **Step 1** In /etc/cinder, create a Huawei Cinder Driver configuration file in .xml format. You can customize the configuration file name, for example, cinder_huawei_conf.xml.
- **Step 2** Configure mandatory parameters in the Cinder Driver configuration file.

```
< RestURL > https://*.*.*:8088/device Manager/rest/; https://*.*.*:8088/device Manager/rest/</RestURL > https://index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.aps//index.a
                 <UserName>***</UserName>
                 <UserPassword>***</UserPassword>
        </Storage>
        <LUN>
                 <StoragePool>***</StoragePool>
                 <LUNType>***</LUNType>
                <LUNCopySpeed>***</LUNCopySpeed>
        </LUN>
        <iSCSI>
                 <Initiator HostName="*****" ALUA="*" FAILOVERMODE="*" SPECIALMODETYPE="*" PATHTYPE="*"</pre>
TargetIP="x.x.x.x y.y.y.y"/>
                 <Initiator HostName="******" ALUA="*" FAILOVERMODE="*" SPECIALMODETYPE="*" PATHTYPE="*"</pre>
TargetIP="x.x.x.x y.y.y.y"/>
     </iSCSI>
</config>
```

Table 4-1 Mandatory parameters

Parameter	Description
Product	Type of a storage product. The value can be V3 , V5 , or Dorado .
Protocol	Type of a connection protocol. Set this parameter to iSCSI for iSCSI scenarios.
RestURL	Access address of the REST interface. If you need to configure multiple RestURL parameters, separate them by semicolons (;).
UserName	User name of a storage administrator.
UserPassword	Password of the storage administrator.
StoragePool	Name of a storage pool to be used. If you need to configure multiple storage pools, separate them by semicolons (;).
LUNType	Type of a LUN. The value can be Thick or Thin . Dorado supports only Thin .
LUNCopySpee d	LUN copy speed. The value can be 1 (low speed), 2 (medium speed), 3 (high speed), or 4 (highest speed).
Initiator	HostName is the regular expression of the Nova or Cinder node host name (For details about the configuration rules and priorities, see the notice). For details about ALUA, FAILOVERMODE, SPECIALMODETYPE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. TargetIP indicates the IP address of the iSCSI target port. You can enter multiple IP addresses and separate them with spaces. This parameter can be set for different hosts.

NOTICE

- 1. Set **Product** based on the actual back-end storage model.
- 2. The value of HostName in the Initiator configuration item is a regular expression. For details about the configuration method, visit **Regular expression**.

When HostName is set to *, the common configuration is used and takes effect on the host with any name. When HostName is set to another value, the general configuration is used.

Configure policy rules:

- Priority: General host name configuration > Common host name configuration For details, see **Example 1** in **7.2 Initiator Configuration Policy Example**.
- Use the first configuration item (for the common configuration of the host name). For details, see Example 2 in 7.2 Initiator Configuration Policy Example.

Step 3 (Optional) Configure optional parameters in the Cinder Driver configuration file.

```
<LUN>
<WriteType>***</WriteType>
<Prefetch Type="***" Value="***"/>
<LUNcopyWaitInterval>***</LUNcopyWaitInterval>
<Timeout>***</Timeout>
<LUNCloneMode>***</LUNCloneMode>
<HyperSyncSpeed>***</HyperSyncSpeed>
<ReplicaSyncSpeed>***</ReplicaSyncSpeed>
<ForceDeleteVolume>***</ForceDeleteVolume>
</LUN>
```

Table 4-2 Optional parameters

Parameter	Default Value	Description
WriteType	1	Cache write type. The value can be 1 (write back) or 2 (write through). For OceanStor Dorado, the value must be 1.
Prefetch Type	3	Cache prefetch policy. The value can be 0 (no prefetch), 1 (fixed prefetch), 2 (variable prefetch), or 3 (intelligent prefetch).
Prefetch Value	0	Cache prefetch value.
LUNcopyWaitInterval	5	After LUN copy is enabled, the plug- in frequently queries the LUN copy progress. You can set a value to specify the query interval.
Timeout	2592000	Timeout interval for waiting the LUN copy of a storage device to complete. The unit is second.

Parameter	Default Value	Description
LUNCloneMode	luncopy	LUN clone mode. The value can be fastclone or luncopy. Only OceanStor Dorado supports fastclone. OceanStor Dorado V300R001C20 and later versions support luncopy. Other storage devices support luncopy only.
HyperSyncSpeed	2	HyperMetro pair synchronization speed. The value can be 1 (low), 2 (medium), 3 (high), or 4 (highest).
ReplicaSyncSpeed	2	Replication pair synchronization speed. The value can be 1 (low), 2 (medium), 3 (high), or 4 (highest).
ForceDeleteVolume	False	When deleting a LUN, if the LUN is in the mapping view, whether to delete it forcibly. Optional values: "True" (forcely), "False" (not forcely)

Step 4 Change the owner and user group of the Huawei Cinder Driver configuration file to be the same as those of the /etc/cinder/cinder.conf file.

-rw-r--r-- 1 cinder cinder 2662 Jul 29 02:13 cinder.conf -rw-r--r-- 1 cinder cinder 778 Jul 30 02:56 cinder_huawei_conf.xml

Step 5 Configure the /etc/cinder/cinder.conf file. Add the following backend configuration to the end of the file:

[huawei_backend] volume_driver = cinder.volume.drivers.huawei.huawei_driver.HuaweiISCSIDriver cinder_huawei_conf_file = /etc/cinder/cinder_huawei_conf.xml volume_backend_name = huawei_backend

■ NOTE

- Backend name huawei_backend can be customized.
- **volume_driver** indicates the type of the driver to be used. Set this parameter to **cinder.volume.drivers.huawei_huawei_driver.HuaweiISCSIDriver**.
- **cinder_huawei_conf_file** indicates the path of the Cinder Driver configuration file.

In the [DEFAULT] area, set enabled_backends to huawei_backend.

[DEFAULT]
...
enabled_backends=huawei_backend

Step 6 If the Cinder Volume service is deployed on multiple nodes, you need to configure a distributed lock in the [coordination] area. For details about <user>, <password>, <host>, and <database>, see the value of connection under [database] in the configuration file.

[coordination]

backend_url = mysql://<user>:<password>@<host>/<database>?charset=utf8

For example, [database] in the configuration file is configured as follows:

```
[database] connection = mysql+pymysql://cinder:302824058e9a4f31@127.0.0.1/cinder
```

Then configure [coordination] as follows:

```
[coordination] backend_url = mysql://cinder:302824058e9a4f31@127.0.0.1/cinder
```

Step 7 Restart the Cinder Volume service.

systemctl restart openstack-cinder-volume.service

----End

4.2 Configuring Cinder Driver for FC-based Storage

This section describes how to configure Cinder Driver for FC-based storage products.

- **Step 1** In /etc/cinder, create a Huawei Cinder Driver configuration file in .xml format. You can customize the configuration file name, for example, cinder_huawei_conf.xml.
- Step 2 Configure mandatory parameters in the Cinder Driver configuration file.

```
<?xml version='1.0' encoding='UTF-8'?>
<confia>
  <Storage>
   <Product>***</Product>
   <Protocol>FC</Protocol>
   <RestURL>https://*.*.*.8088/deviceManager/rest/;https://*.*.*:8088/deviceManager/rest/</RestURL>
   <UserName>***</UserName>
   <UserPassword>***</UserPassword>
 </Storage>
 <LUN>
   <StoragePool>***</StoragePool>
   <LUNType>***</LUNType>
   <LUNCopySpeed>***</LUNCopySpeed>
 </LUN>
  <FC>
   <MinOnlineFCInitiator>xx</MinOnlineFCInitiator>
   <Initiator HostName="xxx" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
   <Initiator HostName="xxx" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
  </FC>
</config>
```

Table 4-3 Mandatory parameters

Parameter	Description
Product	Type of a storage product. The value can be V3 , V5 , or Dorado .
Protocol	Type of a connection protocol. Set this parameter to FC for FC scenarios.
RestURL	Access address of the REST interface. If you need to configure multiple RestURL parameters, separate them by semicolons (;).
UserName	User name of a storage administrator.
UserPassword	Password of the storage administrator.

Parameter	Description
StoragePool	Name of a storage pool to be used.
LUNType	Type of a LUN. The value can be Thick or Thin . Dorado supports only Thin .
LUNCopySpeed	LUN copy speed. The value can be 1 (low speed), 2 (medium speed), 3 (high speed), or 4 (highest speed).
MinOnlineFCInitiator	Minimum available FC initiators. The type is Arabic numerals. The Default value means that all FC initiators are required to be online.
Initiator	HostName is the regular expression of the Nova or Cinder node host name (For details about the configuration rules and priorities, see the notice). For details about ALUA, FAILOVERMODE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. This parameter can be set for different hosts.

NOTICE

- 1. Set **Product** based on the actual back-end storage model.
- 2. The value of HostName in the Initiator configuration item is a regular expression. For details about the configuration method, visit **Regular expression**.

When HostName is set to *, the common configuration is used and takes effect on the host with any name. When HostName is set to another value, the general configuration is used.

Configure policy rules:

- Priority: General host name configuration > Common host name configuration For details, see **Example 3** in **7.2 Initiator Configuration Policy Example**.
- Use the first configuration item (for the common configuration of the host name). For details, see Example 4 in 7.2 Initiator Configuration Policy Example.

Step 3 (Optional) Configure optional parameters in the Cinder Driver configuration file.

```
<LUN>
<WriteType>***</WriteType>
<Prefetch Type="***" Value="***"/>
<LUNcopyWaitInterval>***</LUNcopyWaitInterval>
<Timeout>***</Timeout>
<LUNCloneMode>***</LUNCloneMode>
<HyperSyncSpeed>***</HyperSyncSpeed>
<ReplicaSyncSpeed>***</ReplicaSyncSpeed>
<ForceDeleteVolume>***</ForceDeleteVolume>
</LUN>
```

Table 4-4 Optional parameters

Parameter	Default Value	Description
WriteType	1	Cache write type. The value can be 1 (write back) or 2 (write through).
Prefetch Type	3	Cache prefetch policy. The value can be 0 (no prefetch), 1 (fixed prefetch), 2 (variable prefetch), or 3 (intelligent prefetch).
Prefetch Value	0	Cache prefetch value.
LUNcopyWaitInterval	5	After LUN copy is enabled, the plug-in frequently queries the copy progress. You can set a value to specify the query interval.
Timeout	2592000	Timeout period for waiting LUN copy of an array to complete, in seconds
LUNCloneMode	luncopy	LUN clone mode. The value can be fastclone or luncopy . Only Dorado supports fastclone . Dorado V300R001C20 and later versions support luncopy . Other storage devices support only luncopy .
HyperSyncSpeed	2	HyperMetro pair synchronization speed. The value can be 1 (low speed), 2 (medium speed), 3 (high speed), or 4 (highest speed).
ReplicaSyncSpeed	2	Remote replication pair synchronization speed. The value can be 1 (low speed), 2 (medium speed), 3 (high speed), or 4 (highest speed).
ForceDeleteVolume	False	When deleting a LUN, if the LUN is in the mapping view, whether to delete it forcibly. Optional values: "True" (forcely), "False" (not forcely)

Step 4 Change the owner and user group of the Huawei Cinder Driver configuration file to be the same as those of the /etc/cinder/cinder.conf file.

-rw-r--r-- 1 cinder cinder 2662 Jul 29 02:13 cinder.conf -rw-r--r-- 1 cinder cinder 778 Jul 30 02:56 cinder_huawei_conf.xml

Step 5 Configure the /etc/cinder/cinder.conf file. Add the following backend configuration to the end of the file:

[huawei_backend]
volume_driver = cinder.volume.drivers.huawei.huawei_driver.HuaweiFCDriver
cinder_huawei_conf_file = /etc/cinder/cinder_huawei_conf.xml
volume_backend_name = huawei_backend

□ NOTE

- Backend name huawei_backend can be customized.
- **volume_driver** indicates the type of the driver to be used. Set this parameter to **cinder.volume.drivers.huawei_driver.HuaweiFCDriver**.
- **cinder_huawei_conf_file** indicates the path of the Cinder Driver configuration file.

In the **DEFAULT** area, set **enabled_backends** to **huawei_backend**.

```
[DEFAULT]
...
enabled_backends=huawei_backend
```

Step 6 Restart the Cinder Volume service.

systemctl restart openstack-cinder-volume.service

----End

4.3 Configuring Cinder Driver in Container-based Deployment

This section describes how to configure Cinder Driver for Huawei storage products in container-based deployment. The following uses the kolla deployment method as an example.

Step 1 Configure the **cinder.conf** file in the corresponding path on the host and set related parameters by referring to **4.1 Configuring Cinder Driver for iSCSI-based Storage** or **4.2 Configuring Cinder Driver for FC-based Storage**.

find / -name cinder.conf | grep volume

Step 2 Run the following command to go to the Cinder Volume container:

```
docker ps | grep cinder
docker exec -it -u 0 <cinder_volume_docker_name> bash
```

- Step 3 In the /etc/cinder directory of the Cinder Volume container, create a Huawei Cinder Driver configuration file in .xml format and set related parameters by referring to 4.1 Configuring Cinder Driver for iSCSI-based Storage or 4.2 Configuring Cinder Driver for FC-based Storage. You can customize the configuration file name, for example, cinder_huawei_conf.xml.
- **Step 4** Change the owner and user group of the Huawei Cinder Driver configuration file to be the same as those of the /etc/cinder/cinder.conf file.

```
-rw-----. 1 cinder cinder 2962 Sep 22 07:17 cinder.conf
-rw-----. 1 cinder cinder 443 Sep 14 03:43 cinder_huawei_conf.xml
```

Step 5 Restart the Cinder Volume service.

docker restart <cinder_volume_docker_name>

5 Configuring the Multipathing Software

Multipathing is configured to improve the link reliability of LUNs of SAN storage. If the multipathing configuration is incorrect, an I/O error occurs when a single link is faulty. As a result, the VM file system is read-only or faulty, affecting VM service delivery. When configuring multipathing, install the multipathd service on the host and configure proper parameters. Then, configure multipathing parameters on OpenStack.

5.1 Configuring Host Multipathing

5.2 Configuring OpenStack Multipathing

5.1 Configuring Host Multipathing

5.1.1 Installing the Multipathing Tool Package

Step 1 Install the Multipath tool package based on the operating system.

CentOS:

yum install -y device-mapper-multipath

Ubuntu:

apt-get install -y multipath-tools apt-get install -y multipath-tools-boot

Step 2 Enable the host multipathing service.

CentOS:

/sbin/mpathconf --enable systemctl start multipathd.service systemctl enable multipathd.service systemctl restart multipathd.service

Ubuntu:

systemctl restart multipath-tools.service

5.1.2 Configuring the Multipathing Service

If the AA mode is used, add the following content to the **devices** field in the multipathing configuration file (/etc/multipath.conf). For details, see *OceanStor Dorado V6 Host Connectivity Guide for Red Hat*.

```
devices {
    device {
        vendor "HUAWEI"
        product "XSG1"
        path_grouping_policy multibus
        path_checker tur
        prio const
        path_selector "service-time 0"
        failback immediate
    }
}
```

If the ALUA mode is used, add the following content to the **devices** field in the multipathing configuration file (/etc/multipath.conf). For details, see *Huawei* SAN Storage Host Connectivity Guide for Red Hat.

```
devices {
    device {
        vendor "HUAWEI"
        product "XSG1"
        path_grouping_policy group_by_prio
        path_checker tur
        prio alua
        path_selector "round-robin 0"
        failback immediate
    }
}
```

After the configuration is complete, restart the multipathd service.

systemctl restart multipathd.service

5.2 Configuring OpenStack Multipathing

5.2.1 Configuring Multipathing for OpenStack Nova Nodes

Step 1 Configure the /etc/nova/nova.conf file.

Add volume_use_multipath = True in [libvirt] of /etc/nova/nova.conf.

Ⅲ NOTE

- If "/etc/nova/nova-cpu.conf" exists in the Nova Compute node, add "volume_use_multipath = True" to its [libvirt].
- If the version of OpenStack is Mitaka or earlier, add iscsi_use_multipath = True in [libvirt] of /etc/nova/nova.conf.
- If the version of OpenStack is Newton or later, add volume_use_multipath = True in [libvirt] of /etc/nova/nova.conf.

Step 2 Restart the nova-compute service.

systemctl restart openstack-nova-compute.service

5.2.2 Configuring Multipathing for OpenStack Cinder Nodes

Step 1 Add the following content to the end of BACKEND in the /etc/cinder/cinder.conf file:

[huawei_backend]
...
use_multipath_for_image_xfer = true
enforce_multipath_for_image_xfer = true

Step 2 Restart the cinder-volume service.

systemctl restart openstack-cinder-volume.service

6 Configuring Advanced Properties

Huawei storage supports a series of advanced features. You can create volumes with advanced features by configuring specific volume types.

NOTICE

Licenses for advanced features must be imported onto Huawei storage devices in advance.

- 6.1 Configuring the Thick Property
- 6.2 Configuring the Thin Property
- 6.3 Configuring the SmartQoS Property
- 6.4 Configuring the SmartPartition Property
- 6.5 Configuring the SmartCache Property
- 6.6 Configuring the SmartTier Property
- 6.7 Configuring the HyperMetro Property
- 6.8 Configuring the Remote Replication Property
- 6.9 Creating a LUN in a Specified Storage Pool
- 6.10 Creating a LUN on a Disk of a Specified Type
- 6.11 Creating a LUN in on a Specified Backend
- 6.12 Configuring a Backup Snapshot
- 6.13 Configuring Certificate Authentication
- 6.14 Configuring the Owning Controller of a LUN
- 6.15 Configuring the SmartDedupe Property
- 6.16 Configuring the SmartCompression Property
- 6.17 Creating a Volume Using fastclone

- 6.18 Configuring the LUN Copy Speed
- 6.19 Configuring the Application Type of a LUN
- 6.20 Configuring a Tenant User
- 6.21 Configuring Multiple Backends
- 6.22 Configuring Auto Zoning
- 6.23 Configuring CHAP

6.1 Configuring the Thick Property

This section describes how to configure the Thick property.

NOTICE

OceanStor Dorado does not support Thick volumes.

- **Step 1** Run the following command to create a volume type. *<name>* indicates the name of the volume type.
 - # cinder type-create <name>
- **Step 2** Run the following command to set the Thick volume type. *<vtype>* indicates the volume type name configured in **Step 1**.
 - # cinder type-key <vtype> set capabilities:thick_provisioning_support='<is> true'
- **Step 3** Run the following command to set the scheduling mechanism for the Thick volume. <*vtype>* indicates the volume type name configured in **Step 1**.
 - # cinder type-key <vtype> set provisioning:type='thick'
- **Step 4** Use the volume type in **Step 1** to create a volume.
 - ----End

6.2 Configuring the Thin Property

This section describes how to configure the Thin property.

- **Step 1** Run the following command to create a volume type. <*name>* indicates the name of the volume type.
 - # cinder type-create <name>
- **Step 2** Run the following command to set the Thin volume type. *<vtype>* indicates the volume type name configured in **Step 1**.
 - # cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true'
- **Step 3** Use the volume type in **Step 1** to create a volume.
 - ----End

6.3 Configuring the SmartQoS Property

QoS in OpenStack mainly depends on front-end Hypervisor and back-end storage.

For details about how to configure front-end QoS, see https://docs.openstack.org/cinder/latest/admin/blockstorage-basic-volume-qos.html.

This section describes how to configure back-end QoS for Huawei storage devices.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 Run the following command to enable QoS. *<vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:QoS_support='<is> true'

Step 3 Run the following command to create a QoS specification. <*name>* indicates the name of a customized QoS specification.

cinder qos-create <name> IOType=* <qos_key>=***

Table 6-1 IOType parameter

Parameter	Description	Remarks
ЮТуре	Read/write type.	Valid values are as follows:
		• 0 : read I/O
		• 1: write I/O
		• 2: read and write I/Os

Table 6-2 Optional qos_key parameters

Parameter	Description	Remarks
maxIOPS	Maximum IOPS.	The value is a positive integer, value should not be less than 100 if its type is Dorado.
minIOPS	Minimum IOPS.	The value is a positive integer.
maxBandWidth	Maximum bandwidth.	The value is a positive integer, expressed in MB/s.
minBandWidth	Minimum bandwidth.	The value is a positive integer, expressed in MB/s.
latency	Maximum latency.	The value is a positive integer, expressed in ms.

■ NOTE

- maxBandWidth or maxIOPS cannot coexist with minBandWidth, minIOPS, or latency.
- For OceanStor Dorado, IOType must be set to 2 and minBandWidth, minIOPS, and latency are unavailable.
- **Step 4** Associate the QoS specification with the volume type. <qos_specs> indicates the QoS specification ID created in **Step 3** and <volume_type_id> indicates the volume type ID created in **Step 1**.

cinder gos-associate <gos_specs> <volume_type_id>

Step 5 Use the volume type in **Step 1** to create a volume.

----End

6.4 Configuring the SmartPartition Property

This section describes how to configure the SmartPartition property.

- **Step 1** Run the following command to create a volume type. *<name>* indicates the name of the volume type.
 - # cinder type-create <name>
- **Step 2** Run the following command to enable SmartPartition. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:smartpartition='<is> true'

- **Step 3** Run the following command to configure the name of existing SmartPartition on the back-end storage:
 - # cinder type-key <vtype> set smartpartition:partitionname='***'
- **Step 4** Use the volume type in **Step 1** to create a volume.

----End

6.5 Configuring the SmartCache Property

This section describes how to configure the SmartCache property.

- **Step 1** Run the following command to create a volume type. *<name>* indicates the name of the volume type.
 - # cinder type-create <name>
- **Step 2** Run the following command to enable SmartCache. *<vtype>* indicates the volume type name configured in **Step 1**.
 - # cinder type-key <vtype> set capabilities:smartcache='<is> true'
- **Step 3** Run the following command to configure the name of existing SmartCache on the back-end storage:

cinder type-key <vtype> set smartcache:cachename='***'

Step 4 Use the volume type in **Step 1** to create a volume.

6.6 Configuring the SmartTier Property

This section describes how to configure the SmartTier property.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 Run the following command to enable SmartTier. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:smarttier='<is> true'

Step 3 Run the following command to configure the SmartTier migration policy:

cinder type-key <vtype> set smarttier:policy=***

◯ NOTE

policy indicates a migration policy. The value can be:

- 0: no migration
- 1: automatic migration
- 2: migration to a high-performance tier
- 3:migration to a low-performance tier

Step 4 Use the volume type in **Step 1** to create a volume.

----End

6.7 Configuring the HyperMetro Property

This section describes how to configure the HyperMetro property.

Step 1 Add remote device information to the specified backend in the /etc/cinder/cinder.conf file.

Table 6-3 Parameter description

Parameter	Description	Mandatory
storage_pool	Name of the HyperMetro remote storage pool.	Yes
san_address	Access address of the REST interface.	Yes

Parameter	Description	Mandatory
san_user	User name of a storage administrator or tenant of the HyperMetro remote device.	Yes
san_password	Password of a storage administrator or tenant of the HyperMetro remote device.	Yes
vstore_name	vStore name of the HyperMetro remote device. This parameter needs to be configured only when a tenant user is used.	No
metro_domain	HyperMetro domain name.	Yes
metro_sync_co mpleted	Whether to wait for the synchronization to complete during HyperMetro creation. The value can be True or False . The default value is True .	No
iscsi_info	HostName is the regular expression of the Nova or Cinder node host name (configuration rule and priority description). For details about ALUA, FAILOVERMODE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. TargetIP indicates the IP address of the iSCSI target port. You can enter multiple IP addresses and separate them with spaces. This parameter can be set for different hosts.	Mandatory for iSCSI networking
fc_info	HostName is the regular expression of the Nova or Cinder node host name (configuration rule and priority description). For details about ALUA, FAILOVERMODE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. This parameter can be set for different hosts.	Mandatory for FC networking

- **Step 2** Restart the Cinder Volume service.
- **Step 3** Run the following command to create a volume type. *<name>* indicates the name of the volume type.
 - # cinder type-create <name>
- **Step 4** Run the following command to enable HyperMetro. *<vtype>* indicates the volume type name configured in **Step 3**.
 - # cinder type-key <vtype> set capabilities:hypermetro='<is> true'

Step 5 Use the volume type in **Step 3** to create a volume.

----End

6.8 Configuring the Remote Replication Property

This section describes how to configure the remote replication property.

Step 1 Add remote device information to the specified backend in the **/etc/cinder/cinder.conf** file.

```
replication_device =
backend_id:***,
storage_pool:***,
san_address:https://*.*.*:8088/deviceManager/rest/,
san_user:***,
san_password:******,
vstore_name:***,
iscsi_info: {HostName:xxx;ALUA:xxxx;FAILOVERMODE:xxx;PATHTYPE:xxx;TargetIP:x.x.x.x y.y.y.y};
{HostName:xxx;ALUA:xxx;FAILOVERMODE:xxx;PATHTYPE:xxx;TargetIP:x.x.x.x z.z.z.z},
fc_info: {HostName:xxx;ALUA:xxx;FAILOVERMODE:xxx;PATHTYPE:xxx};
{HostName:xxx;ALUA:xxx;FAILOVERMODE:xxx;PATHTYPE:xxx}}
```

Table 6-4 Parameter description

Parameter	Description	Mandatory
backend_id	Target device ID. This parameter must be specified in the event of failover.	Yes
storage_pool	Storage pool name of the target backend in the event of failover.	Yes
san_address	Access address of the REST interface.	Yes
san_user	User name of a storage administrator or tenant of the remote replication remote device.	Yes
san_password	Password of a storage administrator or tenant of the remote replication remote device.	Yes
vstore_name	vStore name of the remote replication remote device. This parameter needs to be configured only when a tenant user is used.	No

Parameter	Description	Mandatory
iscsi_info	HostName is the regular expression of the Nova or Cinder node host name (configuration rule and priority description). For details about ALUA, FAILOVERMODE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. TargetIP indicates the IP address of the iSCSI target port. You can enter multiple IP addresses and separate them with spaces. This parameter can be set for different hosts.	Mandatory for iSCSI networking
fc_info	HostName is the regular expression of the Nova or Cinder node host name (configuration rule and priority description). For details about ALUA, FAILOVERMODE, and PATHTYP storage ALUA policy, see 9.1 Appendix 1: ALUA Parameter Description. This parameter can be set for different hosts.	Mandatory for FC networking

- **Step 2** Restart the Cinder Volume service.
- **Step 3** Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 4 Run the following command to enable remote replication. *<vtype>* indicates the volume type name configured in **Step 3**.

cinder type-key <vtype> set capabilities:replication_enabled='<is> true'

Step 5 (Optional) Run the following command to set the replication type:

cinder type-key <vtype> set replication_type='<in> ***'

□ NOTE

The value of *** can be:

- **sync**: synchronous replication
- async: asynchronous replication

If **replication_type** is not set, asynchronous replication is used by default.

Step 6 Use the volume type in **Step 3** to create a volume.

----End

6.9 Creating a LUN in a Specified Storage Pool

This section describes how to create a volume in a specified storage pool.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 Run the following command to configure a single storage pool. *<vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set pool_name=StoragePool001

Alternatively, run the following command to configure multiple storage pools:

cinder type-key <vtype> set pool_name="<or> StoragePool001 <or> StoragePool002"

Step 3 Use the volume type in **Step 1** to create a volume.

----End

6.10 Creating a LUN on a Disk of a Specified Type

This section describes how to create a volume on a disk of a specified type.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 Run the following command to configure a disk type. *<vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set disk_type=sas

Alternatively, run the following command to configure multiple disk types:

cinder type-key <vtype> set disk_type="<or> sas <or> ssd"

□ NOTE

The value of disk_type can be ssd, sas, nl_sas, or mix, where mix indicates that two or more of ssd, sas, and nl_sas are used.

Step 3 Use the volume type in **Step 1** to create a volume.

----End

6.11 Creating a LUN in on a Specified Backend

This section describes how to create a volume on a specified backend.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 Run the following command to configure the backend name. *<vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set volume_backend_name=***

□ NOTE

Set **volume_backend_name** to the value of **volume_backend_name** in the **/etc/cinder/cinder.conf** configuration file of the backend.

Step 3 Use the volume type in **Step 1** to create a volume.

6.12 Configuring a Backup Snapshot

This section describes how to configure a backup snapshot.

- **Step 1** Configure a backup driver by referring to https://docs.openstack.org/ocata/config-reference/block-storage/backup-drivers.html. Ensure that the Cinder Backup service is running properly.
- **Step 2** Add the following information to the specified backend in the **/etc/cinder/cinder.conf** file:

```
backup_use_temp_snapshot = True
```

Step 3 Restart the Cinder Volume and Cinder Backup services.

----End

6.13 Configuring Certificate Authentication

This section describes how to configure certificate authentication to connect backend storage.

Step 1 Configure the Cinder Driver configuration file. Add **SSLCertVerify** and **SSLCertPath** in **<Storage>** as follows:

□ NOTE

- **SSLCertVerify** indicates whether to enable certificate authentication. The value can be **True** or **False**. If the parameter is not set, the value is **False** by default.
- SSLCertPath indicates the certificate path for authentication and is valid only when SSLCertVerify is set to True.
- Step 2 Restart the Cinder Volume service.

----End

6.14 Configuring the Owning Controller of a LUN

This section describes how to configure the owning controller of a LUN.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

```
# cinder type-create <name>
```

Step 2 Run the following command to enable the huawei_controller property. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:huawei_controller='<is> true'

Step 3 Run the following command to configure the owning controller name:

cinder type-key <vtype> set huawei_controller:controllername=***

∩ NOTE

The format of the owning controller name is CTEO.A.

Step 4 Use the volume type in **Step 1** to create a volume.

----End

6.15 Configuring the SmartDedupe Property

This section describes how to configure the SmartDedupe property. Only Thin volumes support this property.

Procedure of Enabling SmartDedupe

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 For OpenStack Newton and later versions, run the following command to enable SmartDedupe. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:dedup='<is> true'

For OpenStack Mitaka and earlier versions, run the following command to enable SmartDedupe. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:dedup='<in> true'

Step 3 Use the volume type in **Step 1** to create a volume.

----End

◯ NOTE

OceanStor Dorado supports only Thin volumes. The **capabilities:thin_provisioning_support** parameter in **Step 2** is optional.

Procedure of Disabling SmartDedupe

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 For OpenStack Newton and later versions, run the following command to disable SmartDedupe. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:dedup='<is> false'

For OpenStack Mitaka and earlier versions, run the following command to disable SmartDedupe. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:dedup='<in> false'

Step 3 Use the volume type in **Step 1** to create a volume.

NOTE

OceanStor Dorado supports only Thin volumes. The **capabilities:thin_provisioning_support** parameter in **Step 2** is optional.

6.16 Configuring the SmartCompression Property

This section describes how to configure the SmartCompression property. Only Thin volumes support this property.

Procedure of Enabling SmartCompression

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 For OpenStack Newton and later versions, run the following command to enable SmartCompression. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:compression='<is> true'

For OpenStack Mitaka and earlier versions, run the following command to enable SmartCompression. <vtype> indicates the volume type name configured in Step 1.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true'
capabilities:compression='<in> true'

Step 3 Use the volume type in **Step 1** to create a volume.

----End

■ NOTE

OceanStor Dorado supports only Thin volumes. The **capabilities:thin_provisioning_support** parameter in **Step 2** is optional.

Procedure of Disabling SmartCompression

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

Step 2 For OpenStack Newton and later versions, run the following command to disable SmartCompression. <*vtype>* indicates the volume type name configured in **Step 1**.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:compression='<is> false'

For OpenStack Mitaka and earlier versions, run the following command to disable SmartCompression. <vtype> indicates the volume type name configured in Step 1.

cinder type-key <vtype> set capabilities:thin_provisioning_support='<is> true' capabilities:compression='<in> false'

Step 3 Use the volume type in **Step 1** to create a volume.

□ NOTE

OceanStor Dorado supports only Thin volumes. The **capabilities:thin_provisioning_support** parameter in **Step 2** is optional.

6.17 Creating a Volume Using fastclone

This section describes how to configure **fastclone** to create volumes from snapshots or clone volumes.

Step 1 When creating volumes from snapshots or cloning volumes, specify the -- metadata parameter with the fastclone key as follows:

```
# cinder create --snapshot-id *** --metadata fastclone=True
# cinder create --source-volid *** --metadata fastclone=True
```

----End

◯ NOTE

- If fastclone is not set, LUN copy is used by default.
- Only OceanStor Dorado V3/V6 supports **fastclone**. For details, please contact Huawei storage R&D engineers..

6.18 Configuring the LUN Copy Speed

When you create volumes from snapshots or clone volumes, LUN copy is used by default. This section describes how to configure the LUN copy speed.

When creating volumes from snapshots or cloning volumes, specify the -- **metadata** parameter with the **copyspeed** key as follows:

```
# cinder create --snapshot-id *** --metadata copyspeed=2
# cinder create --source-volid *** --metadata copyspeed=2
```

Table 6-5 copyspeed parameter value description

Parameter Value	Description
1	Low
2	Medium
3	High
4	Highest

Ⅲ NOTE

If **copyspeed** is not set, the medium speed is used by default.

6.19 Configuring the Application Type of a LUN

This section describes how to configure the application type of a LUN.

Step 1 Run the following command to create a volume type. *<name>* indicates the name of the volume type.

cinder type-create <name>

- **Step 2** Run the following command to enable the huawei_application_type property. <*vtype>* indicates the volume type name configured in **Step 1**
 - # cinder type-key <vtype> set capabilities:huawei_application_type='<is> true'
- **Step 3** Run the following command to configure the application type name: # cinder type-key <vtype> set huawei_application_type:applicationname='***'
- **Step 4** Use the volume type in **Step 1** to create a volume.

----End

- This configuration is only available for Dorado V300R001C20 and later versions.
- This configuration includes the configuration of Compression and Dedupe. And its priority is the highest among SmartCompression, SmartDedupe and Application type. Once both configured, Configuration the application type takes precedence.

6.20 Configuring a Tenant User

Configure tenant user parameters in the Cinder Driver configuration file as follows:

Table 6-6 Parameter description

Parameter	Description
UserName	Storage tenant user name.
UserPassword	Storage tenant user password.
vStoreName	vStore name.

NOTICE

- For details about how to configure other parameters, refer to preceding sections.
- Tenant users can be configured only when the storage system version supports the multi-tenant feature.

6.21 Configuring Multiple Backends

Configure multiple storage systems in **/etc/cinder/cinder.conf**. The following is an example.

```
[DEFAULT]
...
enabled_backends = v3_fc,v5_fc
[v3_fc]
volume_driver = cinder.volume.drivers.huawei.huawei_driver.HuaweiFCDriver
cinder_huawei_conf_file = /etc/cinder/cinder_huawei_conf_v3_fc.xml
volume_backend_name = v3_fc
[v5_fc]
volume_driver = cinder.volume.drivers.huawei.huawei_driver.HuaweiFCDriver
cinder_huawei_conf_file = /etc/cinder/cinder_huawei_conf_v5_fc.xml
volume_backend_name = v5_fc
```

6.22 Configuring Auto Zoning

Auto zoning is not compulsory. When you need this function, perform the following configuration.

- **Step 1** Configure the **cinder.conf** file. For details about how to set switch parameters, see https://docs.openstack.org/cinder/latest/configuration/block-storage/fc-zoning.html.
- **Step 2** Restart the Cinder Volume service.

----End

NOTICE

Auto zoning is not supported for tenant users.

6.23 Configuring CHAP

On a public network, any application server whose IP address resides on the same network segment as that of the storage system's iSCSI host port can access the storage system and perform read and write operations. This poses risks to data security. To ensure storage system access security, you can configure CHAP authentication to control application servers' access to the storage system.

Add the following configuration to the Initiator configuration item in the Huawei Driver configuration file:

```
<iSCSI>
    <Initiator ... CHAPinfo="xxx"/>
    <Initiator ... CHAPinfo="xxx"/>
    </iSCSI>
```

Table 6-7 Parameter description

Parameter	Default Value	Description
CHAPinfo	-	User-defined user name and password for CHAP authentication. The format is mmuser;mm-user@storage1. The user name and password are separated by a semicolon (;).
		The CHAP name contains 4 to 223 characters.
		The password contains 12 to 16 characters.
		The password must contain any three types of uppercase letters, lowercase letters, digits, and special characters (` ~!@#\$%^*()=+\ [{}];:,./? and spaces)
		The password must not be the same as the CHAP name or the reverse of the CHAP name.

7 Best Practices

7.1 Configuring Basic Properties for Interconnection Between iSCSI and OceanStor Dorado V3

7.2 Initiator Configuration Policy Example

7.1 Configuring Basic Properties for Interconnection Between iSCSI and OceanStor Dorado V3

The following configurations must be performed after 3 **Deployment** and **5.1 Configuring Host Multipathing** are complete.

Step 1 Configure /etc/cinder/cinder huawei conf.xml.

```
<?xml version='1.0' encoding='UTF-8'?>
<config>
 <Storage>
   <Product>Dorado</Product>
   <Protocol>iSCSI</Protocol>
   <RestURL>https://192.168.128.100:8088/deviceManager/rest/</RestURL>
   <UserName>user</UserName>
   <UserPassword>password</UserPassword>
 </Storage>
 <LUN>
   <StoragePool>StoragePool001</StoragePool>
   <LUNType>Thin</LUNType>
   <LUNCopySpeed>3</LUNCopySpeed>
 </LUN>
 <iSCSI>
   <Initiator HostName="host[0-4]" ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0"</pre>
PATHTYPE="0" TargetIP="128.115.128.100 128.115.128.101"/>
   <Initiator HostName="host[5-9]" ALUA="1" FAILOVERMODE="3" SPECIALMODETYPE="0"</p>
PATHTYPE="0" TargetIP="128.115.128.102 128.115.128.103"/>
 </iSCSI>
</config>
```

Step 2 Modify the owner and owning group of Huawei driver configuration file is the same as those of file /etc/cinder/cinder.conf.

```
-rw-r--r-- 1 cinder cinder 2662 Jul 11 11:25 cinder.conf
-rw-r--r-- 1 cinder cinder 778 Jul 11 11:25 cinder_huawei_conf.xml
```

Step 3 Configure /etc/cinder/cinder.conf.

```
[huawei_backend]
volume_driver = cinder.volume.drivers.huawei.huawei_driver.HuaweiISCSIDriver
```

```
cinder_huawei_conf_file = /etc/cinder/cinder_huawei_conf.xml
volume_backend_name = huawei_backend
use_multipath_for_image_xfer = true
enforce_multipath_for_image_xfer = true
```

Modify the following content in the [DEFAULT] area and add huawei_backend:

```
[DEFAULT]
...
enabled_backends=..., huawei_backend
```

Step 4 Configure multipathing for OpenStack Nova, Add the following content to **[libvirt]** in the **/etc/nova/nova.conf** file

```
[libvirt]
...
volume_use_multipath = True
```

Step 5 Restart the **nova-compute** and **cinder-volume** services.

```
systemctl restart openstack-nova-compute.service
systemctl restart openstack-cinder-volume.service
```

----End

7.2 Initiator Configuration Policy Example

Example 1: (iSCSI) The configuration file content is as follows:

```
<iSCSI>
<Initiator HostName="*" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx" CHAPinfo="xxx"
TargetPortGroup="***"/>
<Initiator HostName="host" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx" CHAPinfo="xxx"
TargetPortGroup="***"/>
</iSCSI>
```

If the host name is host, you can configure ALUA and port groups for the initiator. According to the rule (subject to the precise configuration), the priority of the second record (HostName="host") is higher than that of the first record (HostName="*").

Example 2: (iSCSI)The configuration file content is as follows:

```
<iSCSI>
<Initiator HostName="Host[0-9]" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx" CHAPinfo="xxx" TargetPortGroup="***"/>
<Initiator HostName="Host[5-7]" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx" CHAPinfo="xxx" TargetPortGroup="***"/>
</iSCSI>
```

If the host whose name is **Host6**, the preceding two configuration items meet the configuration rules. According to the rules, the **first** configuration item is used to set the initiator.

Example 3: (FC) The configuration file content is as follows:

```
<FC>
<Initiator HostName="*" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
<Initiator HostName="host" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
</FC>
```

If the host name is host, you can configure ALUA and port groups for the initiator. According to the rule (subject to the precise configuration), the priority of the second record (HostName="host") is higher than that of the first record (HostName="*").

Example 4: (FC) The configuration file content is as follows:

```
<FC>
<Initiator HostName="Host[0-9]" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
<Initiator HostName="Host[5-7]" ALUA="xxx" FAILOVERMODE="XXX" PATHTYPE="xxx"/>
</FC>
```

If the host whose name is **Host6**, the preceding two configuration items meet the configuration rules. According to the rules, the **first** configuration item is used to set the initiator.

8 FAQs

- 8.1 Do I Need to Create Hosts, Host Groups, LUN Groups, and Other Associated Objects on Huawei Storage Before Using Cinder Driver to Map Volumes?
- 8.2 What Can I Do if Mapping Fails Because Nova Logs Indicate that the lsblk Command Fails to Be Executed During Volume Mapping?
- 8.3 Create Volume From Image Failed When verify_glance_signatures=enabled
- 8.4 NewNode1
- 8.5 NewNode2

8.1 Do I Need to Create Hosts, Host Groups, LUN Groups, and Other Associated Objects on Huawei Storage Before Using Cinder Driver to Map Volumes?

No, because Cinder Driver will automatically create associated objects and mapping views. If these objects are manually created on Huawei storage, the mapping logic of Cinder Driver will be adversely affected. Therefore, ensure that these objects are deleted before mapping volumes.

8.2 What Can I Do if Mapping Fails Because Nova Logs Indicate that the lsblk Command Fails to Be Executed During Volume Mapping?

Nova will invoke the **lsblk** command in the volume mapping process to query the information of the mapped volume. This error occurs generally because the lsblk tool is not installed. Ensure that the tool is installed in the Nova system and you have execute permission.

8.3 Create Volume From Image Failed When verify_glance_signatures=enabled

Issue Description

OpenStack executes creating a volume from an image. After copying the image metadata for a period of time, an error is reported: Failed to copy metadata to volume: Glance metadata cannot be updated, key signature_verified exists for volume id b75f957c-0ba4-4099-8706-866c68012779, which eventually causes the volume creation to fail.

Environment

Operating System: CentOS 7.6

Server: Huawei 5288 V5

OpenStack Version: Rocky

Root Cause

BUG Link: https://bugs.launchpad.net/cinder/+bug/1823445

A new field verify_glance_signatures has been added to cinder.image.glance.py since Rocky. The default value is enabled. When CONF.verify_glance_signatures! = 'disabled' and signature_verification in the image also exists, then the DB API does not allow duplicate writes signature_verification flag, thereby throwing an exception.

Solution

In [DEFAULT] of the cinder.conf file, set the verify_glance_signatures field to disabled, for example:

[DEFAULT]

...

verify_glance_signatures = disabled

...

8.4 NewNode1

8.5 NewNode2

9 Appendix

9.1 Appendix 1: ALUA Parameter Description

9.2 Appendix 2: Configuring Volume-backend image

9.3 Appendix 3: Configuring Image-Volume Cache

9.1 Appendix 1: ALUA Parameter Description

ALUA is used to improve the performance of links between hosts and storage. However, the ALUA configuration is complex, involving the storage type, host OS type, and application scenarios (HyperMetro or non-HyperMetro scenarios).

In non-HyperMetro scenarios, see Configuring Multipathing > Configuring Multipathing in Non-HyperMetro Scenarios in the Huawei SAN Storage Host Connectivity Guide for Red Hat.

In HyperMetro scenarios, see Configuring Multipathing > Configuring Multipathing in HyperMetro Scenarios in the Huawei SAN Storage Host Connectivity Guide for Red Hat.

ALUA parameter description

Parameter	Default Value	Description
ALUA	0	0: Huawei UltraPath (default value)1: third-party multipathing software
FAILOVERMODE	-	Initiator switchover mode. This parameter needs to be delivered only when third-party multipathing software is used.
		0: early-version ALUA
		1: common ALUA
		2: ALUA not used
		3: special ALUA

Parameter	Default Value	Description
SPECIALMODETYPE	-	Special mode type of the initiator. This parameter needs to be delivered only when initiator switchover mode is special ALUA. 0: mode 0 1: mode 1 2: mode 2 3: mode 3
PATHTYPE	-	Initiator path type. This parameter needs to be delivered only when third-party multipathing software is used. 0: optimal path 1: non-optimal path

9.2 Appendix 2: Configuring Volume-backend image

This function is provided by the native OpenStack platform. It is used to store images in the storage backend of Cinder to accelerate the speed of creating volumes from images. Official link: https://docs.openstack.org/cinder/stein/admin/blockstorage-volume-backed-image.html

Step 1 Edit the /etc/glance/glance-api.conf file and complete the following configurations:

Configure Cinder as the storage backend of glance and configure the following content in [glance_store]:

```
[glance_store]
...
stores = file,http,swift,cinder
default store = cinder
```

Configure **show_multiple_locations** in [DEFAULT] to parse the locations URL information of Cinder.

```
[DEFAULT]
...
show_multiple_locations = True
```

Step 2 Edit the /etc/cinder/cinder.conf file and complete the following configurations:

Configure **allowed_direct_url_schemes** and **glance_api_version** to use Cinder as the backend of glance.

```
[DEFAULT]
...
glance_api_version=2
allowed_direct_url_schemes = cinder
```

Configure **image_upload_use_cinder_backend** at the backend, for example, [huawei backend].

[huawei_backend]

image_upload_use_cinder_backend = True

Step 3 After the preceding configuration is complete, restart the cinder and glance services.

systemctl restart openstack-glance-* systemctl restart openstack-cinder-*

----End

9.3 Appendix 3: Configuring Image-Volume Cache

This function is provided by the native OpenStack platform. Image volume cache can significantly improve the performance of creating EVS disks from images. Official link: https://docs.openstack.org/cinder/stein/admin/blockstorage-imagevolume-cache.html

Step 1 Obtain the project ID, for example, admin.

openstack project list | grep admin | 366d9f386ecd4defb91df5b0830e7267 | admin |

Step 2 Obtain a Cinder user ID.

openstack user list | grep cinder | 1b2a016cb4a74326a469834dfbc7a841 | cinder

Step 3 Edit the /etc/cinder/cinder.conf file.

Add cinder_internal_tenant_project_id and cinder_internal_tenant_user_id to [DEFAULT].

[DEFAULT]

cinder_internal_tenant_project_id = 366d9f386ecd4defb91df5b0830e7267 cinder internal tenant user id = 1b2a016cb4a74326a469834dfbc7a841

Configure image_volume_cache_enabled at the backend, for example, [huawei_backend].

[huawei_backend]

image volume cache enabled = True

Step 4 After the preceding configuration is complete, restart the cinder and glance services.

systemctl restart openstack-cinder-*

----End